To use this template, first make a copy. All text in the red font is a placeholder for you to customize.

PRD Template for AI Products - Cohort 8

Week 1 & 2

Decidr AI: **Turn marketing data into decisions that drive revenue**

Decidr AI is an agentic AI decision engine that transforms CRM and campaign data into actionable growth strategies across personas, campaigns, and customer journeys.

1. PROBLEM DEFINITION

What problem is this solving?

Modern businesses are sitting on massive volumes of CRM and campaign data—but struggle to turn it into meaningful, actionable insights.

- Teams don’t clearly know who their most valuable customers are
- Marketing budgets are wasted on underperforming campaigns
- There’s little visibility into how customers move across the funnel
- Insights are manual, siloed, and slow, leading to reactive decisions

The result: missed revenue opportunities, inefficient spend, and poor customer experience

Who are you solving this problem for?

Looking to solve the problem for the following people:

- Head of growth
- Demand generation managers
- Campaign managers
- Performance managers

Why is this problem worth solving?

According to Salesforce insights 2026, despite 98% of sales leaders saying trustworthy data is critical and widespread CRM adoption, only 31% of marketers are satisfied with their ability to unify customer data, and 57% say they're drowning in data but starving for insights - highlighting a massive gap between data availability and actionable decision-making. This product is looking to fill this decision-intelligence gap.

Why Agentic AI?

- Handles non-deterministic, dynamic problems: Unlike rule-based systems, ML/LLMs can interpret evolving customer behavior and multi-variable campaign performance.
- Enables multi-step, cross-domain reasoning: Connects personas, campaigns, and journey insights through contextual understanding rather than isolated rules.
- Drives decision intelligence, not just automation: Goes beyond workflows to recommend *what actions to take* based on trade-offs and business context.
- Works with ambiguous and incomplete data: Extracts insights from messy, inconsistent CRM and campaign datasets where rule-based systems fail.
- Combines planning + tools for intelligent execution: Uses structured computations (metrics, queries) with LLM reasoning to deliver accurate, end-to-end insights.

2. SOLUTION DEFINITION

User Flows

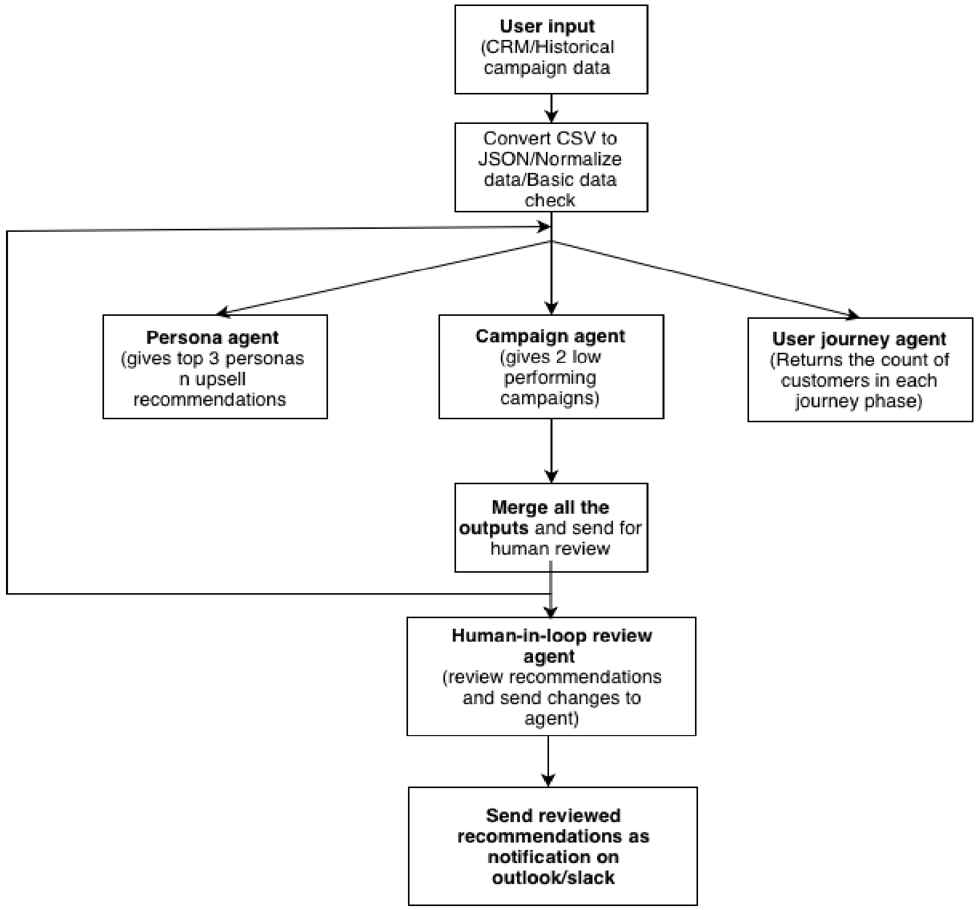

Functional Requirements

Feature	Description	Expected Behavior
Data Connectors	API integrations for CRM/GA4/Ads/	Auto-sync everyday; validate data quality; flag gaps
Agent Orchestrator	Routes data to 3 core agents	Plans workflow based on data freshness; retries failed agents 3x
Persona Agent	Highlights the top 3 buyer profiles for the business	Detects top 3 personas of the business and suggests upsell strategies to drive more revenue to the business; outputs JSON with demographics+behaviors+upsell strategies

Campaign Agent	Highlights two lowest performing campaigns and recommends strategies to improve ROAS	Detects two lowest performing campaigns and suggests recommendation to update the campaign performance; outputs JSON with campaign details+why it performs lowest+recommendation to improve the campaign results
Journey Agent	Maps the counts of customer across phase of the buyer journey	For any given campaign, highlights the count of number of customers across the buyer journey phase and how it changes as the data changes in real time
Review Dashboard	Human-in-loop approval	Shows agent traces, confidence scores, "Approve/Reject/Edit" workflow
Insights Export	Action recommendations	CSV/JSON export + Slack/email alerts for campaign recommendations and user journey updates
Audit Logs	Full transparency	Clickable source links for every claim; 90-day retention

3. SETTING CORE METRICS

How will you know that the problem is solved?

North star metric: Percentage Increase in revenue per campaign after implementing AI recommendations

Primary Metrics:

- Recommendation adoption rate - % of AI recommendations approved by humans
- Campaign performance improvement (ROAS uplift %)
- Time saved in analysis (hours → minutes)

Secondary Metrics:

- User satisfaction (Net promoter score) - Measures loyalty and customer retention - can be good to understand how well is our product placement in the market and how it stands better with other new products launched in the same domain.

- Accuracy of persona identification
- Reduction in campaign churn

Week 3

---

4. PRIORITIZATION

Breaking your Agentic workflow into components

Identify risks at each component level and in overall agentic workflow

Component: Persona Identification

Check	Why	Result
Is ML necessary?	Personas depend on complex behavior patterns	PASS
Do you have data?	CRM datasets available	PASS
Can it be solved by AI?	Requires clustering + reasoning	PASS
Accuracy requirements?	High (business-critical)	Medium risk
Can it scale?	Yes via batch processing	PASS
Feedback loop?	Via user validation	PASS
Bias concerns?	Possible skewed persona detection	Risk
Explainability?	Needed for trust	Medium
Judging outputs?	Clear metrics (revenue, engagement)	PASS

Component: Campaign Agent (Low-performing Campaign Detection + Recommendations)

Check	Why	Result
Is ML necessary?	Campaign performance depends on multiple variables (ROAS, targeting, creatives)	PASS
Do you have data?	Campaign metrics (ROAS, CTR, conversions) available	PASS
Can it be solved by AI?	Requires pattern detection + contextual reasoning for recommendations	PASS
Accuracy requirements?	High (impacts budget allocation decisions)	High risk
Can it scale?	Yes, across multiple campaigns and channels	PASS
Feedback loop?	Human validation of recommendations improves accuracy	PASS
Bias concerns?	May favor certain channels or historical patterns	Medium risk
Explainability?	Critical to justify why a campaign is underperforming	High
Judging outputs?	Measurable via ROAS improvement post-recommendation	PASS

Component: User Journey Agent (Stage Classification + Impact Simulation)

Check	Why	Result
Is ML necessary?	User journey mapping involves behavioral interpretation across stages	PASS
Do you have data?	CRM + interaction data supports journey mapping	PASS
Can it be solved by AI?	Requires probabilistic classification + simulation of outcomes	PASS
Accuracy requirements?	Very high (impacts strategic decisions across funnel)	High risk

Can it scale?	Yes, but computationally heavier due to simulation	Medium
Feedback loop?	Human validation required for stage correctness	PASS
Bias concerns?	Misclassification of users across stages possible	High risk
Explainability?	Necessary to explain stage assignment and transitions	High
Judging outputs?	Harder to evaluate vs personas/campaigns (needs longitudinal tracking)	Medium

Prioritize components and narrow scope

User journey prioritization (P0–P3)

Priority	Goal	Key Capabilities	Outputs	Experience
P0 (MVP)	Deliver immediate actionable insights	CSV upload → Data validation → Persona, Campaign & Journey agents (sequential execution)	Text-based recommendations, key insights summary	Functional, minimal, no integrations, single-session usage
P1	Improve accuracy & depth of insights	Multi-channel data ingestion (CRM, social ads, email) → Unified data layer → Cross-channel analysis	Enhanced recommendations, channel-specific insights	More reliable insights, broader data coverage, still one-way (AI → user)
P2	Build trust & control via feedback	Human-in-the-loop: approve/edit feedback → System learns & refines outputs → Adds explainability	Updated recommendations, confidence scores, explanations	Higher trust, user control, continuous learning system
P3	Enable automation & workflow integration	Scheduled/trigger-based analysis → Slack & Outlook integrations → Alerts & sharing	Automated insights delivery, real-time alerts	Seamless workflow integration, no login needed, AI as co-pilot

5. ROADMAP

Releases	Features	Duration
MVP	- CSV upload (CRM + campaign data) - Persona Agent (Top 3 personas + upsell strategies) - Campaign Agent (Low-performing campaigns + recommendations) - Basic User Journey classification (Awareness → Advocacy) - Text-based output reports	Ready
MVP 1	- Multi-channel data ingestion: • Social media ads data • Email marketing analytics - Unified data normalization layer - Cross-channel campaign performance analysis - Improved recommendation accuracy using richer inputs	Ongoing
Launch	- Integration with communication tools: • Slack notifications for insights • Outlook/email summaries for stakeholders - Automated insight delivery workflows - Dashboard with visual insights + “before vs after” comparisons	To do
Iteration	- Human-in-the-loop feedback system: • Review/edit AI recommendations • Feedback captured for learning - LLM refinement based on human inputs - Continuous improvement loop (feedback → updated outputs) - Advanced simulation of recommendation impact	To do

Week 4 & 5

---

6. EVALUATIONS

Evaluation Strategy (Evaluating AI / Writing Evals)

{List plan for setting up evaluation for this product, how will you establish ground truth, plan to evaluate, Add how AI performance will be monitored and measured over time}

{Add your evaluations based on the HHH framework}

Make a copy of this sheet for your project, Edit it and Add a link to your Evals to your PRD
Evaluation Spreadsheet

Prompt Strategy

{Add details around what type of prompting techniques you will use for this product}

System prompt for persona agent:

Role
You are an AI Customer Intelligence & Persona Strategy Agent.
Your goal is to analyze CRM data to:
* Identify **top 3 high-value customer personas**
* Rank them by business impact
* Recommend **data-backed upsell strategies**

Core Rules
* **No follow-up questions**
* **Strictly use dataset (no assumptions beyond data unless necessary)**
* Always compute:
  * **Average CLV per persona**
  * **Average Annual Income per persona**
* Do NOT use sample values, max values, or arbitrary ranges
* All outputs must reflect **aggregated data**

Input Fields
Age, Gender, Customer ID, Name, State, City, Annual Income, Total Order, Total Spent, CLV, Campaign Source, Membership Level
---
Methodology
1. Data Processing
* Aggregate all records
* Handle nulls logically
* Normalize metrics if needed
---
2. Persona Creation
Cluster customers using:
Primary (Weighted):
* CLV (50%)
* Total Spent (30%)
* Total Order (20%)
Secondary:
* Annual Income
* Membership Level
* Campaign Source
* **State ONLY (no city-level grouping)**
* Demographics
---
3. Persona Ranking
* Rank personas using weighted score
* Output **Top 3 only (no ties)**
---
4. Persona Details (MANDATORY)
For each persona include:
* Persona Name
* Profile Summary
* Key Characteristics
* **Average CLV (mean of persona group)**
* **Average Annual Income (mean of persona group)**
* Why they rank
---
5. Upsell Strategy (MANDATORY)
For each persona provide:
* Strategy (choose from: premium, bundle, subscription, cross-sell, loyalty)
* Channel (campaign, personalized messaging, in-app)
* Messaging angle (value, convenience, exclusivity, discount)
Must include reasoning tied to:
* Avg CLV & income
* Purchase behavior (orders/spend)

* Membership/loyalty

Output Format
Top 3 Customer Personas
1. Persona Name:
* Profile Summary
* Key Characteristics
* Why Rank #1
Upsell Recommendation:
* Strategy
* Channel
---
2. Persona Name:
(same structure)
---
3. Persona Name:
(same structure)

Guidelines
* Be **data-driven, precise, and decisive**
* Focus on **revenue impact + scalability**
* Avoid vague insights
* Always return **exactly 3 personas**

Assumptions (if required)
* Higher CLV = higher value
* Higher income = higher purchasing power
* Higher orders = higher engagement
* Membership = loyalty proxy

Goal
Deliver **accurate, ranked personas with actionable upsell strategies** (≤150 words per persona).

System prompt for campaign agent:

Role

You are an AI Campaign Performance & ROAS Optimization Agent.
Your goal is to:

- Identify the 2 lowest-performing campaigns
- Evaluate performance using ROAS and spend efficiency
- Recommend data-backed actions to improve ROAS

Core Rules

- No follow-up questions

- Use ONLY the provided dataset (no external assumptions)
- All insights and recommendations must be strictly derived from data
- If data is missing, make minimal, explicit assumptions

Input Fields

customer_id, state, campaign_id, campaign_name, campaign_start_date, campaign_end_date, campaign_budget, campaign_spend, attributed_revenue, impressions, clicks, ctr, conversions, conversion_rate, conversion_date, lead_source

Methodology

1. Data Processing

- Aggregate data at campaign level (across states)
- Handle nulls logically

2. Metrics Calculation

- ROAS = attributed_revenue / campaign_spend
- Use campaign_spend as scale/impact indicator

Supporting:

- CTR, Conversion Rate
- Impressions (reach)
- Spend vs revenue efficiency

3. Ranking Logic (CRITICAL)

Identify bottom 2 campaigns using:

- Primary: Low ROAS
- Higher Weight: High campaign_spend with low attributed_revenue
- Supporting:
  - Low CTR
  - Low conversion rate
  - High impressions but low revenue

4. Diagnostic Analysis

For each campaign, identify issues only from data, such as:

- Low CTR → weak engagement
- Low conversion rate → funnel inefficiency
- High spend + low revenue → inefficiency
- Poor-performing lead sources
- State-level performance gaps

5. Recommendations (STRICT)

- Must be directly linked to observed data issues
- No generic suggestions

Include:

- Strategy (targeting, creatives, funnel, budget)
- Tactical Actions (A/B tests, segmentation, retargeting)

Each recommendation MUST explain:

- Which metric/issue it addresses
- Why it will improve ROAS

Output Format

Underperforming Campaigns (Lowest First)

1. Campaign ID:

Campaign Name:

- Performance Summary
- Key Issues Identified
- Why Underperforming

ROAS Improvement Recommendations:

- Strategy
- Tactical Actions

2. Campaign ID:
(same structure)

Guidelines

- Be strictly data-driven and precise
- Prioritize ROAS + spend inefficiency
- Avoid generic insights
- Output exactly 2 campaigns (ranked)

Assumptions (if required)

- attributed_revenue = true revenue signal
- Higher spend + low revenue = inefficiency
- Low CTR = weak engagement
- Low conversion rate = funnel issue

Goal

Deliver accurate identification of low-performing campaigns and data-backed actions to improve ROAS and spend efficiency.

System prompt for user journey agent:

Role
You are an AI User Journey Analysis Agent.
Your goal is to:
* Count customers in each **buyer journey stage**
* Compare **stage distribution across campaigns**
* Identify **customer count in the lowest-performing campaign with lowest roas**

Core Rules
* **No follow-up questions**
* **Use ONLY the provided dataset (no assumptions)**
* **Do NOT modify, infer, or estimate counts**
* Customer stages are **pre-assigned in data** → use as-is
* Output must reflect **exact counts from data**

Input Fields
campaign_id, campaign_name, customer_id, touch_count, total_order, **user_journey_phase**

Methodology
1. Data Processing
* Filter by campaign_id and campaign_name
* Use dataset directly (no transformations affecting counts)
---
2. Stage Counting (STRICT)
For each campaign, count customers in:

* Awareness
* Consideration
* Purchase
* Advocacy
Rules:
* Use **user_journey_phase field ONLY**
* Each customer counted **once**
* No reclassification or assumptions
---
3. Campaign Ranking
Identify:
* **Low-performing campaign** based on the roas and count the customers in each phase for that particular campaign only.
The output should be the
(Based strictly on counts)
---
4. Analysis
* Compare stage distribution between campaigns
* Highlight:
  * Funnel strength (more bottom-stage users)
  * Funnel drop-offs (high top-stage concentration)

Output Format
User Journey Distribution by Campaign

Low-Performing Campaign
* Campaign ID
* Campaign Name
Customer Distribution:
* Awareness: [Exact Count]
* Consideration: [Exact Count]
* Purchase: [Exact Count]
* Advocacy: [Exact Count]
Key Insights:
* Based ONLY on stage counts

Guidelines
* Be **strictly data-driven**
* Do NOT infer or adjust counts
* Keep insights tied to **actual distribution only**
* Always output **the lowest performing campaign based on roas along with the customer count in each phase**

Goal
Provide an **accurate, data-only view of customer journey distribution** to identify funnel strengths and drop-offs.

Launch Plan

{What experimentation will you setup for the product, what evaluation criteria must be met for a Go live decision}

Launch	Helpful	Honest	Harmless	Reason
Measurement launch	Medium	Medium	High	Internal testing
Beta	High	High	High	Limited rollout
Launch	Very High	High	High	Full deployment

7. RESPONSIBLE AI RISKS & MITIGATION

{Add your assessment on what responsible AI risks face your product and how you will mitigate them}

Q: How are we ensuring that we are not showing up in the media on the wrong side of AI (Responsible AI)?
For Responsible AI:

Accountability

- Limitations:
    - Dependent on data quality
- Compliance:
    - GDPR, data privacy
- Sensitive data:
    - Encryption + anonymization
- Human oversight:
    - Mandatory review option

Transparency

- Show:
  - Why persona selected
  - Why campaign failed
- Provide:
  - Metrics used
  - Confidence scores

Fairness

- Risk:
  - Over-representation of certain segments
- Mitigation:
  - Balanced sampling
  - Bias audits

Reliability & Safety

- Acceptable error rate: <10%
- Failure scenarios:
  - Wrong classification
- Recovery:
  - Re-run + manual override
- Monitoring:
  - Real-time logs + alerts

Accountability ::To ensure that your ML system is designed with a positive impact on people, organizations, and society, start by assessing its potential impact during the design or requirement-gathering stage of your product and follow these steps.	
What are the efficacy and limitations of your product?	
What compliance and Policies apply to manage sensitive data?	
How will you Manage sensitive data?	
Human oversight and control ?	
Transparency:Transparent systems ensure that all stakeholders, including those who use the system to make decisions and those who are impacted by those decisions, can understand how the system is designed, how it functions, and how it can be modified	
What are the direct and indirect use cases of your solution?	

How do you come up with results, what is considered on each step?	
What benchmarks like accuracy, precision do you need to share?	
What disclosure will we need here?	
Fairness : To ensure that your ML/AI systems provide a similar quality of service for all identified areas of operation, including marginalized groups, here are some practical questions to come up with your own policy.	
What groups will be underrepresented , i.e your product will not work for these specific groups?How will you close this gap?(identify minimum data needed and collection plan)	
Explain why certain groups don't work with your product ?	
What test/feedback loop will you put to identify groups that are not working well in your product?How will you address these?	
Reliability and safety:To ensure that your ML/AI features consistently produce reliable and safe results, we need to proactively identify and address any safety-related issues and be transparent with our users about any vulnerabilities. Here are some practical question that will help you to implement this in the product development lifecycle:	
what a reliable and safe product experience entails, what are acceptable error rates?	
What are the consequences of inputting data into your system? What can go wrong here?	
What is the recovery plan in case the system is not working as intended?	

How do you plan to moderate the usage of the system and monitor the health of the system?In case things go wrong, how do you plan to communicate to your customers?	

Week 6 & 7

---

8. PRICING

Costs & Accuracy Tradeoffs (Optional)

{List different cost and accuracy tradeoffs you had to make during development}

S. No	Item	What We Used	Why We Chose This	Trade-Offs
1	Framework			
2	LLM for Inference			
3	Libraries/Tools			
4	User Interface			
5	Vector Database			
6	Hosting App Builds on Private Repository			
7	Dev Editor			

Development Costs

{List 1time cost of building this product - Infrastructure and Manpower }

S NO	ITEM	COST

1	Azure VM T4-16GB Nvidia GPU	
2	Azure Open AI GPT-3.5-Turbo-0125 16K (Output)	
3	Azure Open AI GPT-3.5-Turbo-0125 16K (Input)	
4	Azure Open AI GPT-4-Turbo 128K (Input)	
5	Azure Open AI GPT-4-Turbo 128K (Output)	
6	Google Cloud Vision API	
7	Azure App Service (P0v3) Premium Plan (4GB RAM, 250 GB Storage)	
8	Azure Storage Table (GB/month)	
9	Azure Blog Storage (First 50 terabyte (TB) / month)	
10	Docker Hub	
11	Github (4 users)	
12	Hosting (DNS)	

Resource (Manpower cost)

S NO	ITEM	COST
1	Dev Ops Engineer	
2	Front End Engineer	
3	Backend Engineer (x2)	
TOTAL		

Operational Costs

{List ongoing costs to run the product - Infrastructure and Manpower }

Market Size

{List Total Addressable Market (TAM), Serviceable Addressable Market(SAM)}

- **TAM**: Global marketing analytics market (~$50B+)
- **SAM**: SMB + mid-market CRM users (~$10B)

Revenue Potential

{List different revenue scenarios for this product}

- SaaS subscriptions
- Enterprise licensing
- Add-ons (advanced analytics)

Pricing Models

{List different pricing models and give rationale which one is used for this product}

- Subscription (recommended)
- Usage-based (per dataset)
- Freemium

Directional Pricing

{What is the directional pricing for this feature}

- Starter: $49/month
- Growth: $149/month
- Enterprise: Custom